



Security Analysis of BOINC

Karl Chen, Paul Huang

Department of Computer Science – University of California, Berkeley

email: {quarl, pbhuang}@cs.berkeley.edu



Abstract

Berkeley Open Infrastructure for Network Computing (BOINC) is a software platform for distributed-computing using volunteered computer resources. It generalizes the software-engineering aspects of SETI@home-type projects; it's the "@home" in SETI@home. We analyze why the security defenses in BOINC are of paramount importance, discuss possible attacks on BOINC, and propose and analyze defenses. We also compare some other distributed-computing products.

Motivation

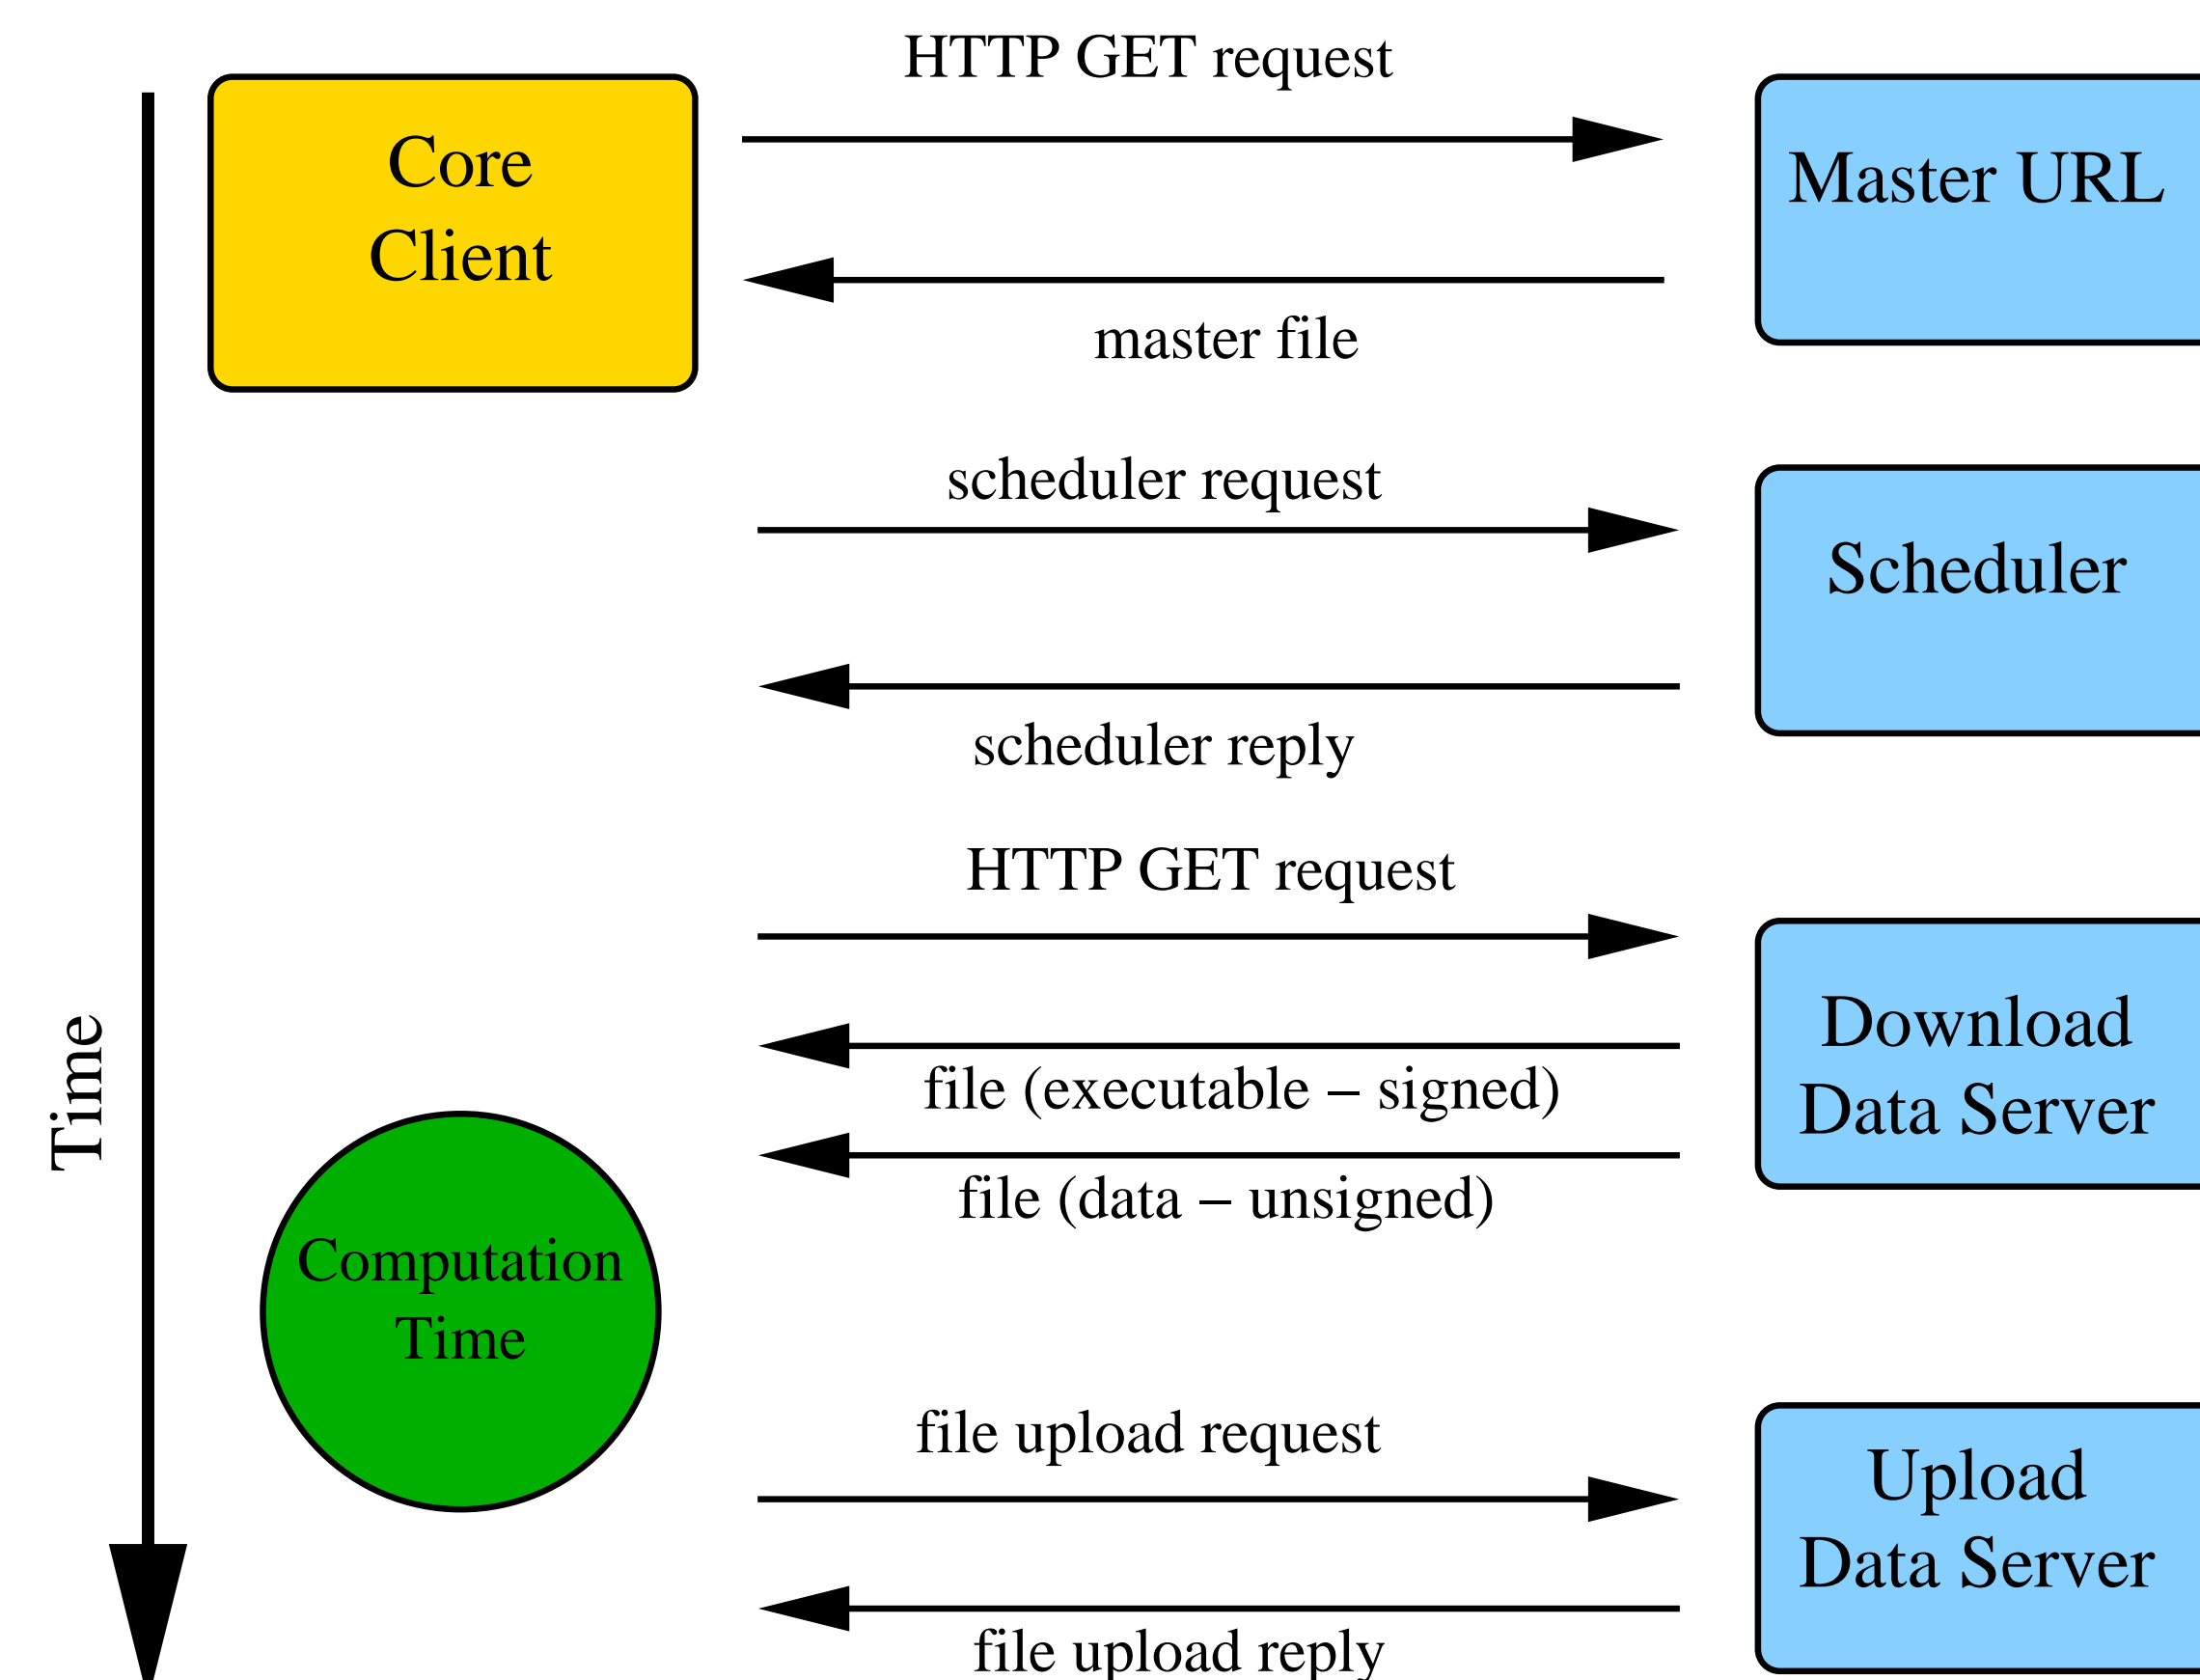
Why is compromise of a BOINC server more dangerous than compromise of www.mozilla.org?

- 4 million users running the client software (soon)
- Automatic, periodic download and execution of binary applications; default interval of 1 day
- Compromised server \approx malicious server
- BOINC clients assume non-malicious servers
- Compromise unnoticed for 1 day \Rightarrow all clients compromised
- Mass-compromises of clients (stealing credit card numbers, deleting files) can lead to the end of public distributed computing

Worms possible

- Attack on server + attack on client \Rightarrow worm!
 - Clients are 99% Windows/x86 monoculture
 - Servers are Linux/Solaris running Apache + MySQL + PHP + BOINC
- Hop over firewalls, infect non-public servers
- Servers also know clients' IP addresses & OS versions

BOINC Architecture Overview

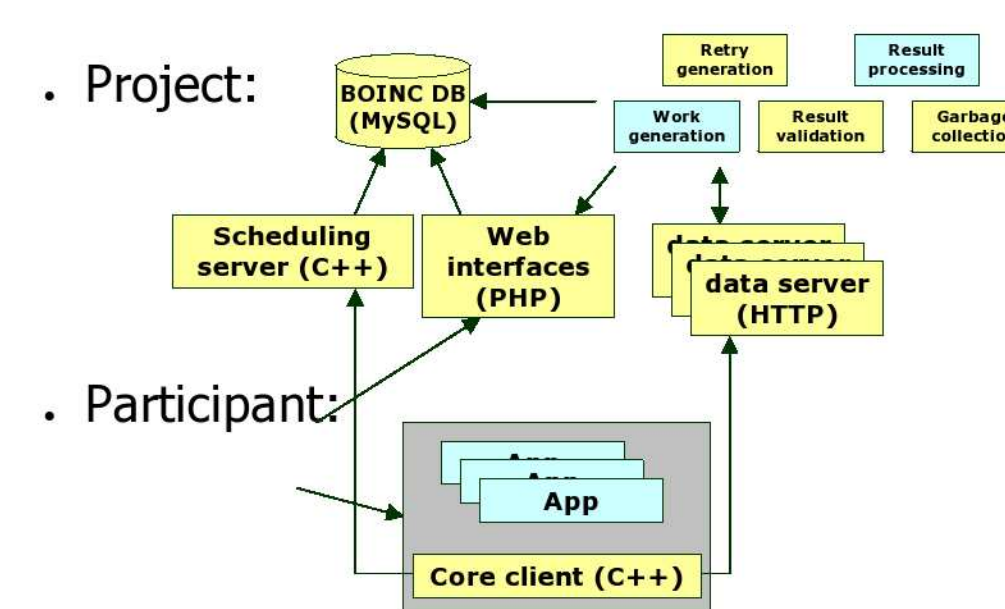


Server components

- Master URL
- Schedulers
- Download Servers
- Upload Servers

Non-user-visible:

- Database, daemons



Threat Model

- 1 million computers could be compromised in 1 day
- Attackers may invest high resources
 - Man-in-the-middle attacks
 - Find holes in any services running on BOINC servers

Current Defenses

- Binary executables are signed using RSA
- Server components can be separated
- Uploading files requires signature (to prevent DoS on upload server)

Insufficient!

General problem of running untrusted applications; signatures not the best solution. If private key compromised:

- Keys don't have lifetime (violates fail-safe defaults)
- Re-keying protocol: sign new key with old key
- Only creates a race between project managers and attacker to re-key - and attacker has a head start

Not only executables can be exploited (violates principle of complete mediation):

- Applications use non-type-safe languages (so far)
- Applications invoke external programs
- Complexity \Rightarrow security verification difficult
- Applications in the past have had security holes
- Attack on input files \Rightarrow arbitrary code execution
 - "tar": attacker could pass "--rsh-command" if not careful
 - "tar x", embedded zip library: no check for "." / "/" when extracting
 - * Overwrite arbitrary files, including executables

Attacks on input files:

- Client follows pointers from Master URL \rightarrow Schedulers \rightarrow Download servers
- Compromise anywhere in path \Rightarrow attack on input files
 - Vulnerable to DNS attacks, MITM attacks
 - Master URL contains embedded XML list of schedulers, as well as "user of the day" profile, news, etc. (violates the principle of least common mechanism)

Analysis of Other Architectures

Some other architectures:

- MoneyBee
- Classic Folding@Home
- UnitedDevices
- Distributed.net
- D2OL

Other architectures:

- Most lack feature of downloading executables (less insecure)
- Most have less levels of indirection to download (less insecure)
- Most download unsigned input files (equally insecure)
- Some rely on secret key embedded in application (insecure)
- Many rely on security through obscurity (insecure)
- One uses Java (very secure)

GRID architectures:

- Servers and clients generally assume each other trusted, so little security

Defenses

Specific issues:

- tar, Zip library: add checks for paths
- Master URL: don't re-use "bells-and-whistles" page
- Compromise of input files: sign them also
- Projects can choose to use type-safe language

Treat application as untrusted:

- Sandbox each application: feasible, partial solution (chroot)
- Interpreted language: not feasible in short term (legacy applications)
- Virtual machine: not feasible (performance)
- Proof-carrying code: feasible, but need it for external programs as well (one-time compiler cost)
- Intercept and mediate syscalls: feasible (Janus for Solaris, Linux)